

Assignment 1

Submission Instructions

1. Please submit the assignment 1 on submission page linked from the course webpage.
2. Upload the source code files (Matlab/Python) for the programming assignment as a single compressed (zip) file. Follow the code structure provided in code-data.zip (Matlab implementation). The code for problem 4 should be in a folder named "Problem-4" and for problem 5 in folder "Problem-5". Please create subfolders for each sub-question with name as the sub-question index (as given in code zip).
3. You should submit part of your assignment(Problem 4 and Problem 5) as a soft copy(pdf) generated using LaTeX. All graphs should be generated programmatically and should be included in the report with proper numbering. Solutions to Problem 1,2,3 can be submitted as a soft copy (along with problem 4 and problem 5)or hard copy(handwritten). It is highly encourage to submit the whole solution as a single pdf report.
4. The handwritten solutions for problems 1, 2, 3 can be submitted directly to the TAs in their office mentioned in course webpage.
5. Use the pdf submission link in the assignment submission page to submit your pdf.
6. Upon multiple submissions till the deadline, the most recent submission will overwrite the previous one.
7. Please download and use the code blocks provided in the assignment submission page for doing the programming assignments. Please follow the folder structure/template provided. Follow the same folder structure as given in the Matlab template, for Python implementations too.
8. Please make sure that the file formats are correct, i.e. code is in zip format and report is in pdf format. TA's will not evaluate defaulting assignment submissions.
9. Datasets are included in the data folder of respective problems in the hw1_code_data.zip file
10. Questions about the assignment can be posted on the group email if or directed to the TAs {annervaz.km, aadirupa.saha} @csa.iisc.ernet.in, sharmistha@grad.cds.iisc.ac.in.

Some useful concepts and definitions

1. Training Set and Test Set: A typical machine learning model is trained using some data. The subset of data used for learning model parameters is called Training Set. Test set of the subset of data (it's intersection with train data is a null set), used to evaluate your model on unseen data. Good performance on the test set indicates generalisation capability of the model.
2. Cross Validation: Typically, the data is divided into two groups, train and test. The evaluations obtained in this case tend to reflect the particular way the data are divided up. The solution is to divide up data in such a way that, we test the model on each data point atleast once. This method is called cross-validation. Following is a sketch of the procedure.
 - (a) Divide data randomly into k folds (subsets) of equal size.

- (b) Train the model on $k-1$ folds, use remaining one fold for testing.
 - (c) Repeat this process k times so that all folds are used atleast once for testing.
 - (d) Compute the average test performance on these k test sets.
3. Precision, Recall, Accuracy: Important metrics to evaluate the performance of learned algorithms. Read more here: https://www.cs.cornell.edu/courses/cs578/2003fa/performance_measures.pdf
4. One-Hot-Encoding: One way to convert text into a numerical representation in machine learning algorithms is to treat each word as a feature. One such representation is called one-hot-vector encoding. In this representation, each document (text sample) is represented as a vector of n dimension, where n is the total number of unique words appearing in the entire text corpus. Each index in the vector corresponds to the count of the the vocabulary word in the document.

Questions

1. **Probability Review.** Suppose that X and Y are discrete random variables. Their joint *pmf* function $P(X,Y)$ is given by the table below. Show how to compute $P(X)$, $P(Y)$, $P(X/Y)$ and $P(Y/X)$ (use conditional probabilities). [5 Mark]

		X		
Y		x1	x2	x3
	y1	0.1	0.2	0.1
	y2	0.1	0.2	0.3

2. Given a 2-class classification problem in 2 dimensional space. What is the Bayes decision boundary for this problem. [5 Mark]

$$p(x|w_1) \sim N(0, I)$$

$$p(x|w_2) \sim N([1, 1]^T, I)$$

$$\text{and } P(w_1) = P(w_2) = 0.5$$

3. [10 Mark] Consider a cost-sensitive binary classification task in which misclassifying a positive instance as negative incurs a cost of $c \in (0, 1)$, and misclassifying a negative instance as positive incurs a cost of $1 - c$. This can be formulated as a learning task with instance space \mathcal{X} , label and prediction spaces $\mathcal{Y}, \hat{\mathcal{Y}} \in \{\pm 1\}$, and cost-sensitive loss $\ell_c : \{\pm 1\} \times \{\pm 1\} \rightarrow \{0, c, 1 - c\}$ defined as

$$\ell_c(y, \hat{y}) = \begin{cases} c & \text{if } y = -1 \text{ and } \hat{y} = 1 \\ 1 - c & \text{if } y = 1 \text{ and } \hat{y} = -1 \\ 0 & \text{if } \hat{y} = y. \end{cases}$$

Let D be a joint probability distribution on $\mathcal{X} \times \{\pm 1\}$, with marginal distribution μ on \mathcal{X} and conditional label probabilities given by $\eta(x) = \mathbf{P}(y = 1|x)$, and for any classifier $h : \mathcal{X} \rightarrow \{\pm 1\}$, define

$$\text{er}_D^c[h] = \mathbf{E}_{(x,y) \sim D} [\ell_c(y, h(x))].$$

Derive a Bayes optimal classifier in this setting, i.e. a classifier $h^* : \mathcal{X} \rightarrow \{\pm 1\}$ with

$$\text{er}_D^c[h^*] = \inf_{h: \mathcal{X} \rightarrow \{\pm 1\}} \text{er}_D^c[h].$$

How does your derivation change if the loss incurred on misclassifying a positive instance as negative is a and that for misclassifying a negative instance as positive is b for some arbitrary $a, b > 0$?

4. **Programming Exercise: Logistic Regression.** [30 Mark] Write a piece of MATLAB code to implement logistic regression on a given binary classification data set (\mathbf{X}, \mathbf{y}) , where \mathbf{X} is an $m \times d$ matrix (m instances, each of dimension d) and \mathbf{y} is an m -dimensional vector (with $y_i \in \{0, 1\}$ being a binary label associated with the i -th instance in \mathbf{X}): your program should take the training data (\mathbf{X}, \mathbf{y}) as input and output a d -dimensional weight vector \mathbf{w} and bias (threshold) term b representing the learnt classifier. For this problem, you are provided a spam classification data set¹, where each instance is an email message represented by 57 features and is associated with a binary label indicating whether the email is spam or non-spam. The last column of 'spambase.data' denotes whether the e-mail was considered spam (1) or not (0). The data set is divided into training and test sets. The goal is to learn from the training set a classifier that can classify new email messages in the test set.

- (a) Use your implementation of logistic regression to learn a classifier from 10% of the training data, then 20% of the training data, then 30% and so on upto 100%. In each case, measure the classification error on the training examples used, as well

¹UCI Machine Learning Repository: <http://archive.ics.uci.edu/ml/datasets/Spambase>.

as the error on the given test set (you can use the provided code `classification_error.m` to help you compute the errors). Plot a curve showing both the training error and the test error (on the y -axis) as a function of the number of training examples used (on the x -axis). Such a plot is often called a **learning curve**. (Note: the training error should be calculated only on the subset of examples used for training, not on all the training examples available in the given data set.)

- (b) Incorporate a L_2 -regularizer in your implementation of logistic regression (taking the regularization parameter λ as an additional input). Run the regularized version of logistic regression on the entire training set for different values of λ in the range $\{10^{-7}, 10^{-6}, 10^{-5}, 10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 1\}$. Plot the training and test error achieved by each value of λ (λ on the x -axis and the classification error on the y -axis) and identify the value of λ that achieves the lowest test error (resolving ties in favor of the smallest value of λ). Now, select λ from the same range through 5-fold cross-validation on the training set (the train and test data for each fold are provided in a separate folder `Problem-4\data\spambase-cross-validation`); report the average cross-validation error (across the five folds) for each value of λ and identify the value of λ that achieves the lowest average cross-validation error (again resolving ties in favor of the smallest value of λ). Does the cross-validation procedure select the right value of λ ?

Hints: The `glmfit` function in MATLAB, which implements (unregularized) logistic regression, can be used for sanity check in this problem; note that the input labels to this function must be in $\{0, 1\}$ and not in $\{\pm 1\}$. For implementing your own version of logistic regression, the `fminunc` MATLAB function for unconstrained optimization will be useful. Also, do not forget to include the bias (threshold) term in the logistic regression model.

5. Programming Exercise: Naive Bayes, Perceptron and Winnow Algorithm. [50 Mark]

A dataset for sentiments of movie review is available at <http://ai.stanford.edu/~amaas/data/sentiment/>. The dataset contains movie reviews in natural language and their sentiment whether they are positive[1] or negative[-1]. We have processed this large dataset into three train sets [small, medium and large] and one test set. Use train-small dataset for training in sub-question a-f. Use all three train sets for question part g. There is a single test set, which should be used for all sub-questions. The dataset contains processed documents, the last column in each line denotes the label [1,-1], all other columns of vector represent encoded data [14666 columns]. These are large files and may take about 5 mins to load in matlab running on a desktop/laptop. Each index on the vector corresponds to the count of a vocabulary word [one-hot-vector encoding], `imdb_vocab.csv` file contains the word for index i at line i (use this information for sub-question d).

Code for accuracy prediction is included in the code folder (`Classification_Accuracy.m`). You can use 'confusionmat' function for computing confusion matrix in matlab (`sklearn.metrics.confusion_matrix` in python).

- Implement a naive bayes classifier for document classification for sentiment analysis dataset. Please report train and test accuracies. Also provide the confusion matrix for test results.
- Implement Perceptron algorithm for the same problem of document sentiment classification. Run your code for the first 10000 training samples for 25 rounds (A round is an iteration over the training data). After each round report the training and test accuracies in a plot, with round number on x-axis, train & test accuracies on y-axis.
- Implement Winnow algorithm for the same problem of document sentiment classification. Run your code for the first 10000 training samples for 25 rounds. After each round report the training and test accuracies in a plot, with round number on x-axis, train & test accuracies on y-axis.
- After complete training in the perceptron algorithm, list top 10 words, ordered based on the magnitude of w_i is the final weight vector w for perceptron. For example, if index $i=60$ is one of the top 10 w_i , then word at line 60 is one of the influential features for this task prediction. *Non-credit: Why do you think these words have high coefficient value in the final weight vector? Hint: relate to the problem of sentiment classification*
- Shuffle your training dataset using given shuffle script (`shuffle.m`), train a perceptron using the shuffled data and note the number of updates that you are making to the weight vector. Repeat shuffle and train cycle 20 times. Give a histogram plot of the number of updates you made for the 20 shuffles of the training dataset. What can you say about γ from these values?*

- (f) In the perceptron training algorithm you were making an update to the weight vector when making a mistake as $w + (y * x)$. Now change this update rule as $w + (\eta * y * x)$, where η is a small value, typically known as learning rate. Train your perceptron using different values of η ranging from 0.05 to 0.6 with increments of 0.05. For each value report the number of updates made, train and test accuracies in a table. Plot the number of updates (y-axis) made against η on x-axis.
- (g) Run-time comparison: Train perceptron algorithm using train-small, train-medium, train-large datasets and test on given test dataset. State the training time for 3 the train sets in a plot, with number of examples in train set on x-axis and runtime on y-axis. Report test accuracy on models trained on these 3 train datasets in a table. State your conclusion about the run-time complexity of perceptron.